

Scalable Architecture

Scalability - Not a metric by itself

- Ability of a system to ensure that all other -ilities are either enhanced or maintained, and, not adversely impacted, when load / volume increases!
- System & Development Scalability

Path Towards MSA Decompose into independent, differentially treatable and smaller units!


Scaling First Decompose Functionally (Y-axis), then scale each service by x or z axis

Modularity Basics: Coupling & Cohesion, SRP, Composition vs Aggregation; Avoid Shotgun Surgery & Divergent Change

Y-Axis Scaling – Questions to ask: Actors, Usecases, Entities, Volume, Resource Usage, Security, Criticality, Team Structure

Software Architecture Basics

- It's all about decomposition
- 4+1 Architecture View: Logical, Implementation, Process & Deployment Views.
- Styles: Layered (Prez, Biz, Data) – Hexagonal (Dependency Inversion)



Monolith

- Can be perceived at any level: Function, Class, Module, App, Suite of Apps
- Scaling: X & Z axis of Scale Cube
- Pros: Low Latency, Low Maintenance (Build, Deployment, Config), Simple to Scale
- Cons: High Load Time, Hard to maintain/replatform or to scale/treat differentially

MSA Challenges High Latency – Multiple Points of Failures – Data Consistency Issues – Operational Challenges – Decomposition is a challenge that requires great skills

Popular MSA Patterns Synch RPC (REST, gRPC), Asynch (Messaging Q), Circuit Breaker, Service Discovery Patterns, API Gateway, Saga (compensating transactions), Deployment Patterns (Mesh, VM, Container, Serverless)

4+1 Arch View	Monolith	Microservices
Implementation View	Single Component	Set of multiple components
Logical View	No specification	No specification, (mostly hexagonal)

	SOA	MSA
Scope	Enterprise	Local
Granularity	Coarse-grained	Fine-grained
Communication	Mostly Synchronous	Mostly Asynchronous

Scalable Database

	Replication	Sharding
Why?	Availability - Scale out – Low Latency	Scale Out – Fault Isolation – Performance
Challenges	Synchronization, Failure Handling	Fair Distribution

Distributed Consistency – Replication Methods – Network Faults are inevitable - Leads to CAP Theorem – When distributed, P is not a choice, so it is between CP & AP - ACID → CP, BASE → AP

Safety Guarantees given by a transaction are described by ACID

Atomicity: Either All or None – Commit or Rollback/Abort – doesn't deal with Concurrency

Isolation: Concurrency - Serializability – 2 Phase Locking – Optimistic Concurrency Control

Consistency: Validation Rules

Durability: Guaranteed replication and write to Storage from Volatile Memory

BASE Basically Available – Soft State – Eventually Consistent

RDBMS	ACID	CA (CP is a challenge)
NoSQL	ACID or BASE	CP or AP



Firewall: Intercept & Filter
Forward Proxy: Client Anonymity
Reverse Proxy: Server Anonymity
App Router: Decide & Redirect
Load Balancer: Load Distribution

LB Algos: Round Robin, Weighted Round Robin, Least Connections, Least Response Time
LB Types: Layer4 (TCP) & Layer7 (HTTP – App Awareness)

Reverse Proxy
Security: Authentication, DDoS Prevention
Performance: Caching, SSL Termination, Compression

Session Management

Session Identifier Options: IP Addr, Browser Login, URL Rewriting, Hidden Form Fields, Cookies
Distributed env: Sticky Sessions – DB to store – Distributed Data Store (like Redis) – No Server Side Sessions (JWT)

Performance responsiveness of a system to execute any action within a given time interval

Latency Time spent waiting for other action(s) to be completed

Availability Percentage of time the service remains operational under normal circumstances in order to serve its intended purpose. $A = \frac{\text{Total time} - \text{Sum of downtime}}{\text{Total time}}$

Performance Considerations


- Offload View logic to Client
- No server side sessions
- Separate core from less critical use cases
- Choose the right communication protocol
- Leverage CDN & Caching
- Separate Real-time, near real-time and batch work
- Prefer Async Messaging Queue, wherever applicable – increases reliability too.

Availability Patterns: Failover & Replication – Solution to Single Point of Failures

Replication: Master-Slave, Master-Master (either CP or AP)

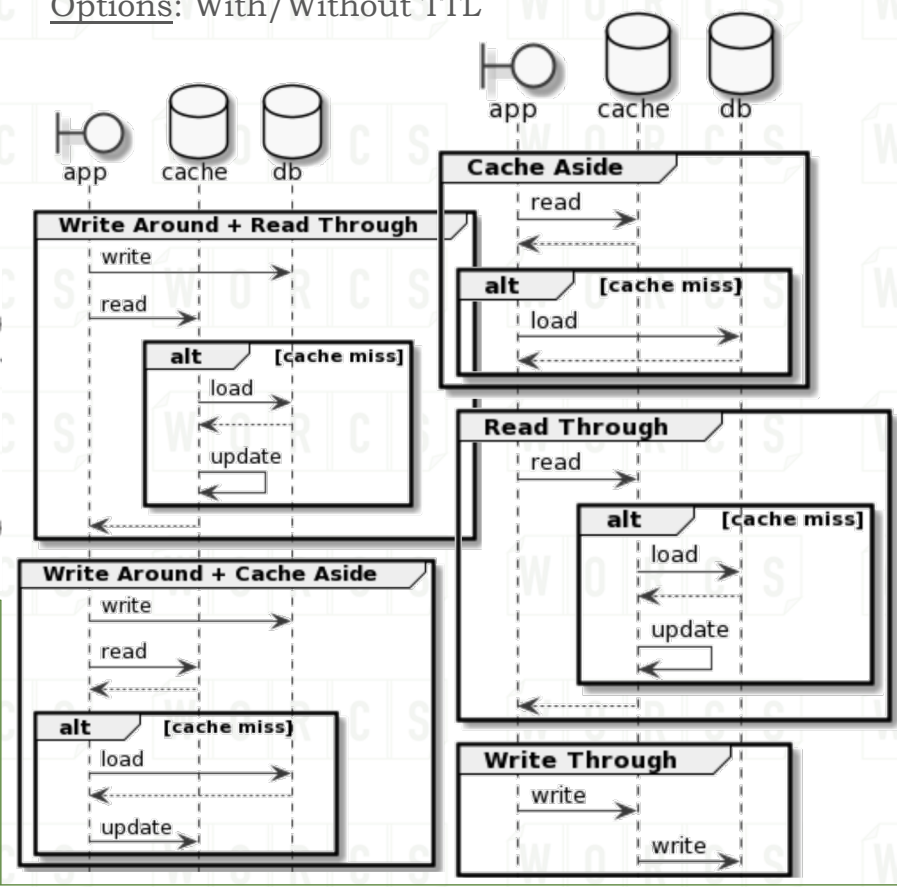
Caching

- Client Caching (End User)
- Edge Caching (closer to end users)
- Web Server Caching (Reverse Proxies, etc.)
- Application Caching
- Database Caching (DB Buffers, Indexing, etc.)



DB Caching Strategies:

Cache Aside, Read Through, Write Through, Write Around, Write Back/Behind, Refresh Ahead
Options: With/Without TTL



Capacity Planning:

Determining the production capacity required to meet changing demands – Periodic Task – Key Resources: CPU, Network, Memory, Disk – Measurement is better than intuition

Capacity Planning Steps

1. Determine Requirements: Prod Owner / Biz Analyst - User Types, Common/Frequent Use Cases, Short/Long Term Growth – Average/Peak Traffic, Expected Response Time
2. Run Baseline Tests: Test env ~ prod – mock 3rd party calls – stress test common cases for long duration
3. Extrapolate/Estimate: Use the metrics & TPS to extrapolate